

The Re-ID Problem with an Animal Dataset

Team 10

Connor Kannally
Sal Hargis





Problem Introduction and Motivation: Animal Re-ID

- Re-ID is re-occurrence of an individual animal (have you seen *this* animal before?)
- Challenging problem, individuals change over time, each individual has different viewpoints
- Species classification not sufficient for ecology studies



AnimalCLEF25 Dataset

- Multi-Species Individual Animal Identification
- More difficult to identify individuals than species:
- 10k individuals and 140k images
- Variety of animals including sea, land, and air animals
- Many different viewpoints and backgrounds (close-up or far away)



Same
individual
“Nannan”

Evaluation Method

Closed Set:

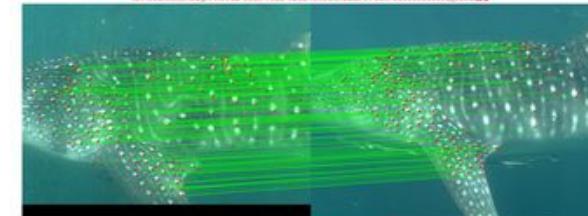
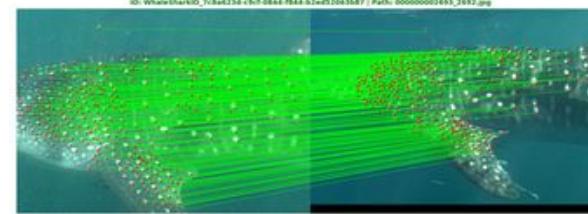
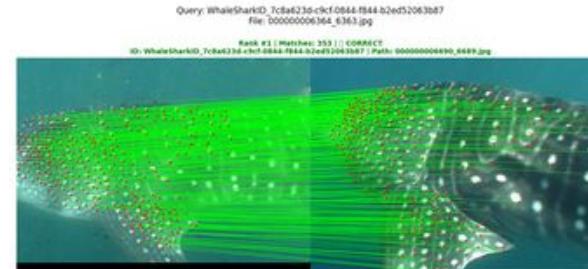
- Compare overall accuracy of baseline method (Resnet50) to baseline + LightGlue (feature matching)
- R-1 and R-20 Accuracy

Open set recognition accuracy with:

- 2000 images of individuals
- 500 images not in training set

LightGlue Feature Matching with Whale Shark

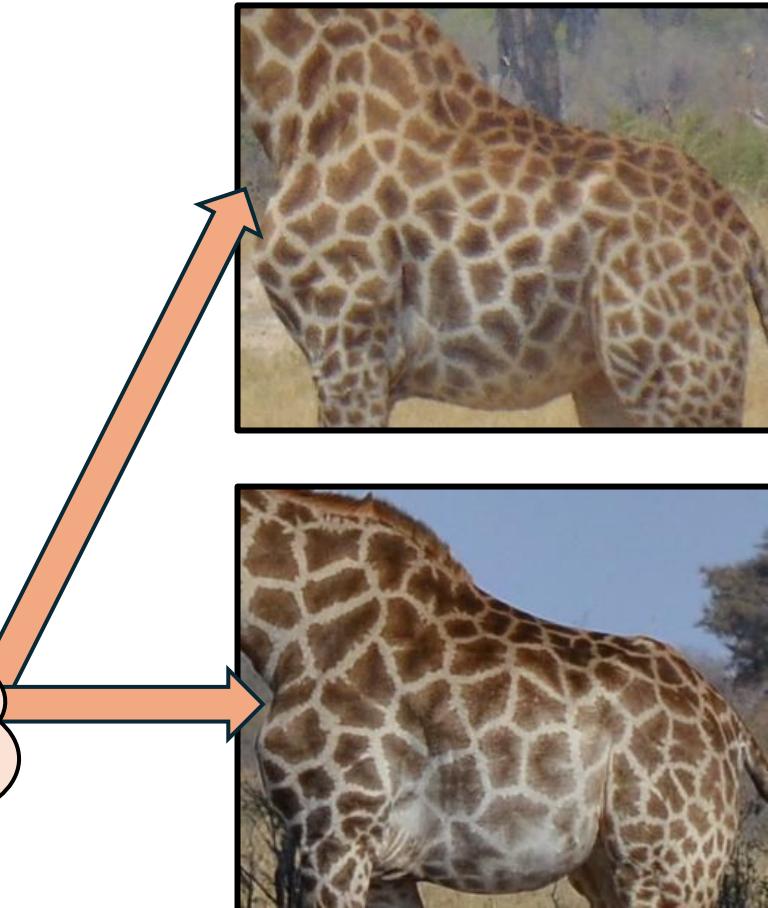
Open set recognition tests individuals NOT in the training set



Baseline Approach

- Fine-Tuning out of the box ResNet50 to the Wildlife-10k dataset
- Baseline Resnet50 on AnimalCLEF25:
 - ~99% accuracy for species identification (near perfect – wrong task)
 - ~56% for individual identification (poor performance – correct task)

ResNet50 output
these as Giraffes, but
not as ***different***
Giraffes!



Advanced Approach: Model Pipeline

Multispecies Animal Re-ID Using a Large Community-Curated Dataset

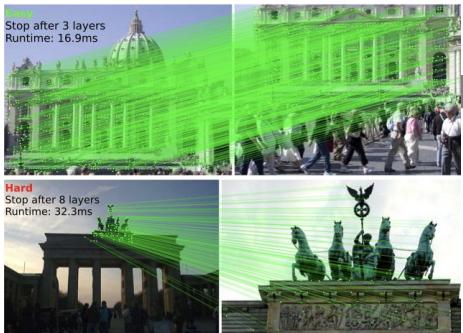
Lasha Otarashvili^{*}
Conservation X Labs

Tamilselvan Subramanian
Conservation X Labs

Jason Holmberg
Conservation X Labs

J.J. Levenson
Bureau of Ocean Energy Management: Washington, DC, US

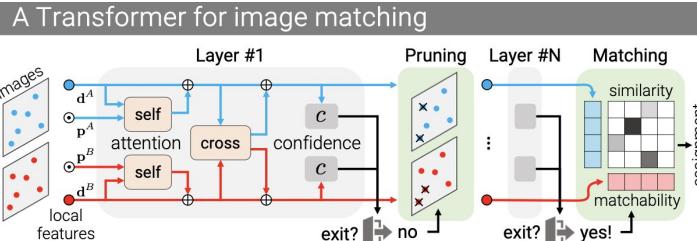
Charles V. Stewart
Rensselaer Polytechnic Institute



Coarse-Grain: Top 100 candidates for each individual using:

- ResNet50
- EfficientNet-B3

Fine-Grain Ranking:
Re-Rank the top 100 with Light Glue



Advanced Approach: ArcFace Loss for Comparisons

- Optimized with ArcFace Loss instead of Softmax
 - Softmax only cares about separability, no mechanism to group related images together (important for Re-ID!)
 - ArcFace creates an embedding space where Angle = Similarity
- ResNet50 with ArcFace Loss
 - Embedding size: 512 dim
 - Standard baseline
- EfficientNet-B3 with ArcFace Loss
 - Embedding size: 512 dim
 - Optimized architecture using Wildlife-10k

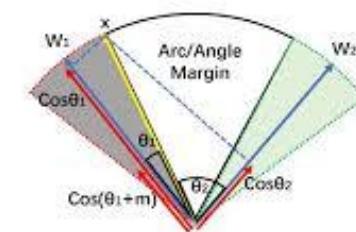


Figure 2. Geometrical interpretation of ArcFace. Different colour areas represent feature spaces from distinct classes. ArcFace can not only compress the feature regions but also correspond to the geodesic distance on the hypersphere surface.

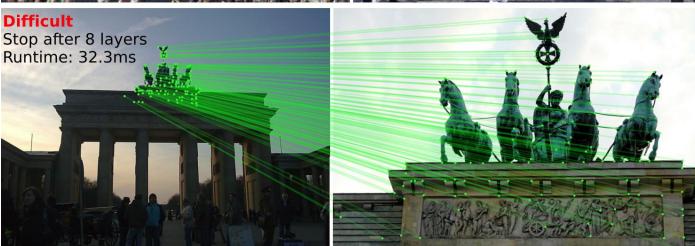
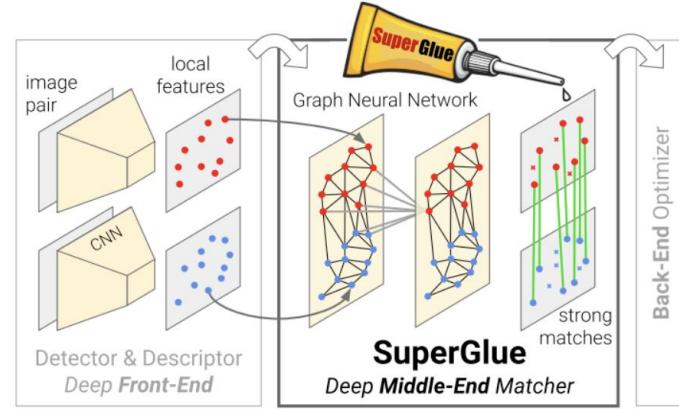
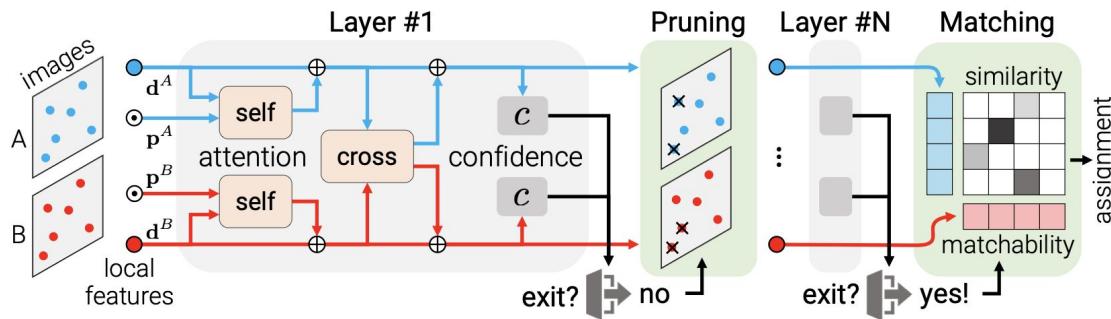
<https://arxiv.org/pdf/1801.07698>

Advanced Approach: SuperGlue

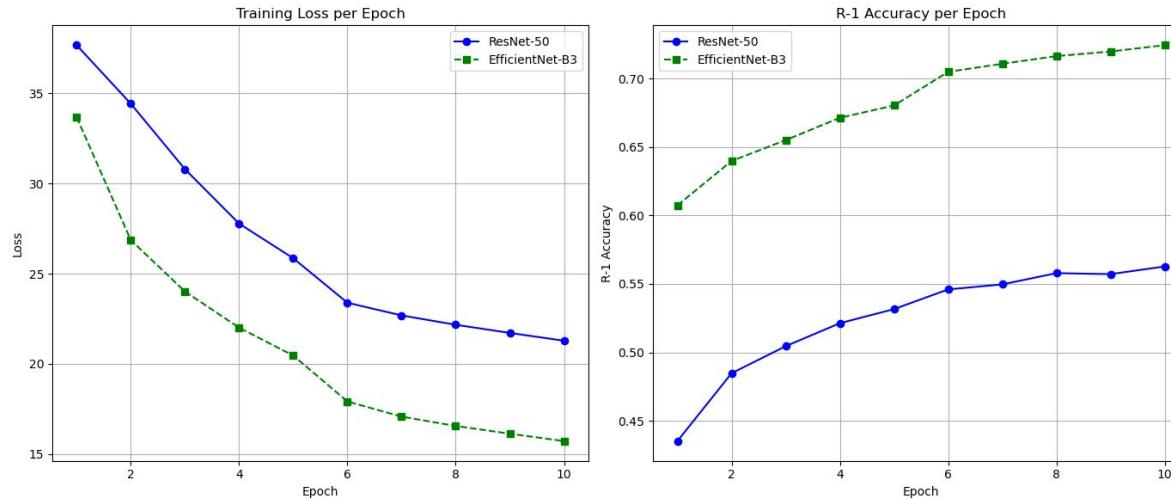
Motivation: Models like ResNet fail when the viewpoint changes (strongly impacts the image vector)

SuperGlue: Can verify identity based on strict physical markings

LightGlue: Stops early on “easy” matches to be faster



Results and Analysis: Baseline Approach (without LightGlue)



Why did EfficientNetV3 go down?

- LightGlue could be looking at the background

Takeaways:

- EfficientNetV3 Accuracy higher alone than with LightGlue

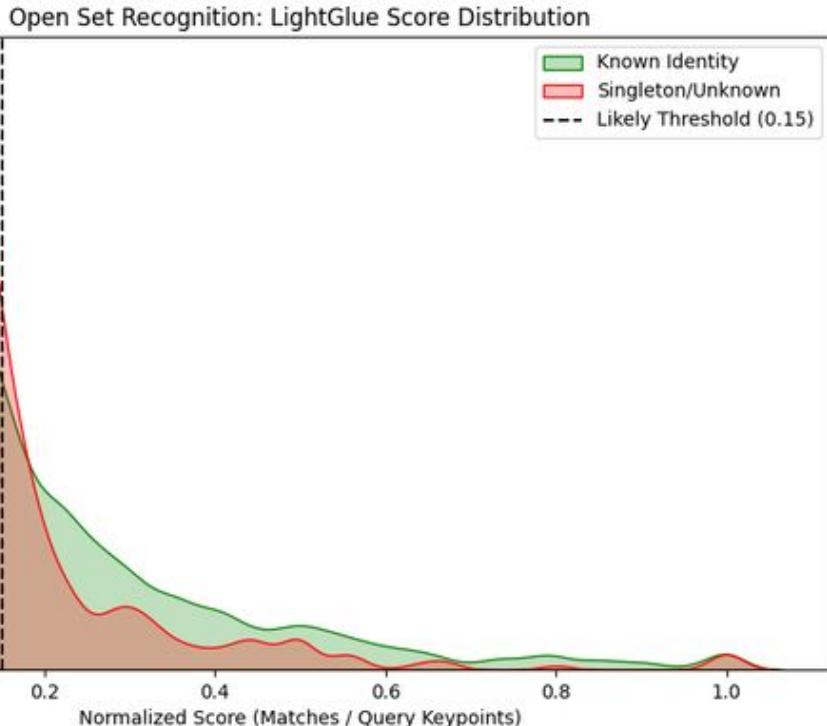
ResNet50 vs. EfficientNetB3

- R-1 Accuracy
 - ResNet50: ~56%
 - EfficientNetV3: ~72% (highlighted)
- R-20 Accuracy:
 - ResNet50: ~76%
 - EfficientNetV3: ~86%

ResNet50 vs. EfficientNetB3 (with LightGlue)

- R-1 Accuracy
 - ResNet50: ~64%
 - EfficientNetV3: ~71% (highlighted)
- R-25 Accuracy:
 - ResNet50: ~79%
 - EfficientNetV3: ~85%

Results and Analysis: LightGlue Challenge for Re-ID



Feature Matching and Re-ID:

- Where to set the threshold for a new individual
 - No way to set a threshold to be able to distinguish between known and new individuals

Takeaway:

- LightGlue can help re-rank when accuracy of baseline model is lower
- Not sufficient for identifying new individuals (given our context)

Workload

1-2 weeks setting up with large amounts of data prepared

3-4 weeks exploring various approaches
(SuperGlue, LightGlue, Resnet50,
EfficientNetV3, ViT, etc)

1 week of evaluation and analysis of approaches



Challenges



**Downloading and setting up OSC
remotely**



**Having a plan of path forward
before taking time and resources
to train models (time consuming)**



**Getting a good candidate images
for LightGlue (without taking too
long)**



Future Work

- Explore different feature matching algorithms for re-ranking
- Explore different loss functions other than Triplet and ArcFace
- **Data Preprocessing:**
 - Remove background, so the model learns only features of the individual (image segmentation)
 - Data augmentation to generate alternative and flipped views of individuals
 - A way to label the images based on view (left, right, top, etc)

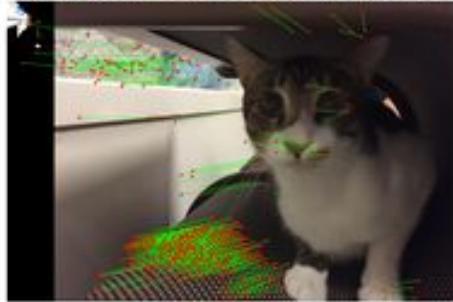
*Feature matching focusing
only on the background*

Query Path: /fs/scratch/PAS3163/Kennedy_Hangs/unified_reid_data/CatIndividualImages_183/0145_0144_0145_044.jpg
True ID: 2399

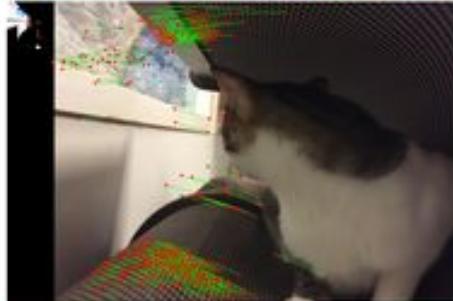
Rank #1 | Matches: 319 | : CORRECT
ID: 3286 | Path: /fs/scratch/PAS3163/Kennedy_Hangs/unified_reid_data/CatIndividualImages_183/0143_0140_0140.jpg



Rank #2 | Matches: 250 | : CORRECT
ID: 3286 | Path: /fs/scratch/PAS3163/Kennedy_Hangs/unified_reid_data/CatIndividualImages_183/0143_0141_0145_045.jpg



Rank #3 | Matches: 219 | : CORRECT
ID: 3286 | Path: /fs/scratch/PAS3163/Kennedy_Hangs/unified_reid_data/CatIndividualImages_183/0143_0140_0149_049.jpg



Thank You!